



Proposed-system advantage in low-to-moderate-load regimes

S=32, budget=8 GPUs



S=64, budget=12 GPUs



Dynamic scheduling and pruning are compared at the same fixed GPU budget.

—●— All-HBF colocated —▲— Static hetero + prune —★— Proposed dynamic + prune