


Why $\mu=2$ wins at 64K context, S=128, budget=16
Dynamic EP with c=.98; A:M optimized independently



μ_1 is serial (TBT=A+M); $\mu \geq 2$ uses the steady-state overlap bound TBT=max(A,M).