

Per-model decode latency — TPOT p99 (H3 vs SPaKV)

