


Sustainable decode throughput per GPU under SLO (context = 64K)

TPOT SLO = 50 ms



TPOT SLO = 100 ms

