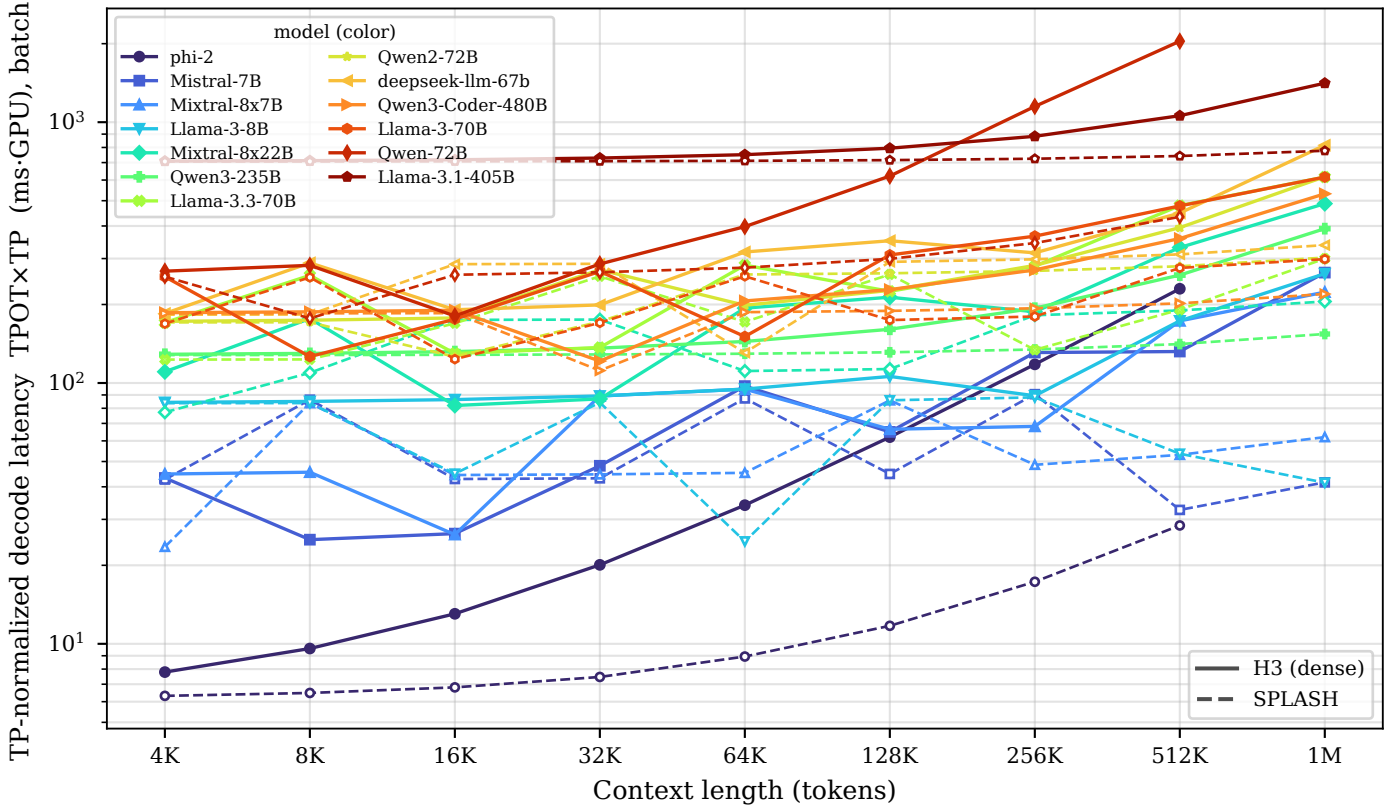
Decode latency per token, TP-normalized: H3 (solid) vs SPLASH (dashed) — Blackwell
model (color)
phi-2
Mistral-7B
Mixtral-8x7B
Llama-3-8B
Mixtral-8x22B
Qwen3-235B
Llama-3.3-70B
Qwen2-72B
deepseek-llm-67b
Qwen3-Coder-480B
Llama-3-70B
Qwen-72B
Llama-3.1-405B
H3 (dense)
SPLASH
TP-normalized decode latency TPOT×TP (ms·GPU), batch 1
10^3
10^2
10^1
4K
8K
16K
32K
64K
128K
256K
512K
1M
Context length (tokens)