


(a) Same 10% token budget: quality **(b) Same 10% token budget: performance**



(c) Attention-level latency



(d) Why page granularity is faster

